

UiT

THE ARCTIC
UNIVERSITY
OF NORWAY

Mandatory Assignment 3

INF-2200 (fall 2016)

Department of Computer Science

University of Tromsø



In a sentence

Implement a hierarchical memory sub-system with two cache levels, used by a CPU simulator.

Details

- Two cache levels
 - Level 1
 - Read-only instruction cache
 - Data cache
 - Level 2
 - Unified cache
- Precode provides CPU simulator that will perform memory accesses against your memory subsystem.
 - You will implement `memory.c`
 - No changes to precode necessary!

Goal

- Best cache design
 - Use benchmark from assignment 1 (or other benchmark).
 - Measure cache hit and miss ratio.
 - Experiment with different parameters.

Precode

- Small and straight forward.
- Implements API and starting point for memory subsystem.
- Implements CPU simulator that does
 - Instruction fetch
 - Load data
 - Store data
- Memory trace stored in binary format
 - P2addrTr struct (byurt.h)
- Use valgrind to generate memory trace logfile.
- Precode provides script for converting logfile to binary trace file read by CPU simulator.

Generate trace

- Step 1
 - Run the following command:
`valgrind --log-file=logfile --tool=lackey --trace-mem=yes [your-program-name]`
 - This will create a file `trace.tr` that contains the memory trace of your program.
- Step 2
 - Parse the trace file by running:
`python traceconverter.py`
 - This will produce a file `logfile` that can be used as input to the cache simulator.
- Step 3
 - Run the cache simulator:
`./cachesim logfile`
 - The precode will initialize your memory subsystem by calling `memory_init()` and will then, for each memory access in the `logfile`, call one of the functions:
 - `memory_fetch()` – if the memory access is an instruction fetch.
 - `memory_read()` – if the memory access is a data read.
 - `memory_write()` – if the memory access is a data write.

Requirements

- Implementable in real hardware
 - Parameters realistic according to book
- Parameters
 - Easily changeable
 - Start with parameters given in assignment text
- L1 data cache and L2 unified cache should support both reads and writes.
 - L1 instruction cache should be read only.

Requirements

- Simplifications
 - Assume each data access is within boundary of one cache line.
 - Assume all instructions of fixed size, and aligned.
 - Neither is actually true on x86...
 - Not necessary to implement reads and writes that actually *transfer data*, just count cache hits and misses.
- Select a replacement policy and implement it
 - Random
 - LRU
 - Temporal/Spatial
- Write policy
 - Write-back
 - Write-through

Count

- Hits and misses
 - Differentiate
 - Layers
 - Reads
 - Writes

Method of approach

- Evaluate cache by creating memory trace
 - Use valgrind and convert with python script
- Two traces
 - Correctness (hits and misses known in advance)
 - May have do be created manually
 - Trace from benchmark

Report

- Cache performance tweaks
 - How?
 - Why? / Why not?
 - Temporal/Spatial
- Correctness test
- Reductions and simplifications

Deadline

- November 2nd @ 12:00 PM (noon)
- Hard deadline (no extensions possible)!